Crunchy Container Overview

Crunchy Data Solutions, Inc.

Version 1.2, 2016-04-20

crunchy-postgres

The crunchy-postgres container executes the Postgres database.

Packages

The container image is built using either the Crunchy Postgres release or the community version based upon a flag in the Makefile.

The Crunchy postgres RPMs are available to Crunchy customers only. The Crunchy release is meant for customers that require enterprise level support.

The PGDG community RPMs can be used as well by simply commenting out the Crunchy yum repo within the Dockerfiles and uncommenting the PGDG yum repo.

setup.sql

The setup.sql script is used to define startup SQL commands that are executed when the database is first created.

Environment Variables

- PG_MODE - either master or slave, this value determines whether the database is set up as a master or slave instance
- PG_MASTER_USER - the value to use for the user ID created as master. The **master** user has super user priviledges.
- PG_MASTER_PASSWORD - the password for the PG_MASTER_USER database user
- PG_USER - the value to use for the user ID created as a normal user. This user is created as part of the setup.sql script upon database creation and allows users to predefine an application user.
- PG_PASSWORD - the password for the PG_USER database user that is created
- PG_DATABASE - a database that is created upon database initialization
- PG_ROOT_PASSWORD - the postgres user password set up upon database initialization

Features

The following features are supported by the crunchy-postgres container:

- use of openshift secrets
- ability to restore from a database backup
- use of custom pg_hba.conf and postgresql.conf files
- ability to override postgresql.conf configuration parameters

Example

Examples of using the container are found in the standalone and openshift documents.

crunchy-backup

The crunchy-backup container executes a pg_basebackup against another database container. The backup is a full backup using the standard utility included with postgres, pg_basebackup.

Backup Location

Backups are stored in a mounted backup volume location, using the database host name as a subdirectory, then followed by a unique backup directory based upon a date/timestamp. It is left to the user to perform database backup archives in this current version of the container. This backup location is referenced when performing a database restore.

Dependencies

The container is meant to be using a NFS or similar network file system to persist database backups.

Example

Examples of using the container are found in the standalone and openshift documents.

crunchy-collect

Description

Postgresql metrics are collected by the crunchy-collect container. To start collection of metrics on a Postgres database, you add the crunchy-collect container into the pod that holds the crunchy-pg container.

Requirements

Metrics are stored in the crunchy-scope container. crunchy-scope runs the prometheus time series database, the prometheus pushgateway, and the grafana web application. An example of starting the crunchy-scope container is in examples/openshift/scope.json.

To start the crunchy-scope container, run the following:

```
oc process -f scope.json | oc create -f -
```

The crunchy-scope data in this example is stored in emptyDir volume types. To persist the data and

grafana templates long term, you will want to use NFS volume types as specified in examples/openshift/scope-nfs.json.

When running crunchy-scope, the following ports are available: * crunchy-scope:9090 - the prometheus web user interface * crunchy-scope:9091 - the prometheus pushgateway REST API * crunchy-scope:3000 - the grafana web user interface

crunchy-collect Environment Variables

- POLL_INT - number of minutes to sleep until metrics are collected. defaults to 15 minutes
- PROM_GATEWAY - the http URL of the prometheus pushgateway into which the metrics will be pushed. defaults to http://crunchy-scope:9091

Example

An example of using crunchy-collect is stored in examples/openshift/master-collect.json

This example collects metrics for the pg-master database. Run the example as follows:

```
oc process -f master-collect.json | oc create -f -
```

crunchy-pgbadger

The crunchy-pgbadger container executes the pgbadger utility. A small http server is running on the container, when a request is made to:

```
http://<<ip address>>:10000/api/badgergenerate
```

Environment Variables

- BADGER_TARGET - only used in standalone mode to specify the name of the container, also used to find the location of the database log files in /pgdata/$BADGER_TARGET/pg_log/*.log

Features

The following features are supported by the crunchy-pgbadger container:

- basic invocation of pgbadger against the database log files

Example

Examples of using the container are found in the standalone and openshift documents.

crunchy-pgpool

The crunchy-pgpool container executes the pgpool utility. Pgpool can be used to provide a smart postgres-aware proxy to a postgres cluster, both master and slave, so that applications can only have to work with a single database connection.

Postgres slaves are read-only whereas a master is both read and write capable.

Environment Variables

- PG_USERNAME - user to connect to postgres
- PG_PASSWORD - user password to connect to postgres
- PG_MASTER_SERVICE_NAME - database host to connect to for the master node
- PG_SLAVE_SERVICE_NAME - database host to connect to for the slave node

Features

The following features are supported by the crunchy-pgpool container:

- basic invocation of pgpool

Example

Examples of using the container are found in the standalone and openshift documents.

crunchy-watch

We create a container, crunchy-watch, that lives inside a pod with a master container.

The watch container has access to a service account that is used inside the container to issue commands to openshift.

You set up the SA using this:

oc create -f sa.json

You then set up permissions for the SA to edit stuff in the openshift project, this example allows all service accounts to edit resources in the 'openshift' project:

~~~~~ oc policy add-role-to-group edit system:serviceaccounts -n openshift ~~~~~

You then reference the SA within the POD spec:

watch-logic.sh is an example of what can run inside the watch container

I copy the oc command into the container from the host when the container image is built. I could not find a way to install this using the redhat rpms, so I manually add it to the container.

Environment Variables

- SLEEP_TIME - the time to sleep in seconds between checking on the master
- PG_MASTER_SERVICE - the master service name
- PG_SLAVE_SERVICE - the slave service name
- PG_MASTER_PORT - database port to use when checking the database
- PG_MASTER_USER - database user account to use when checking the database using pg_isready utility
- PG_DATABASE - database to use when checking the database using pg_isready

Logic

The watch container will watch the master, if the master dies, then the watcher will:

- create the trigger file on the slave that will become the new master
- change the labels on the slave to be those of the master
- will start watching the new master in case that falls over next

Example

Examples of using the container are found in the standalone and openshift documents.

crunchy-dba

The crunchy-dba container implements a cron scheduler. The purpose of the crunchy-dba container is to offer a way to perform simple DBA tasks that occur on some form of schedule such as backup jobs or running a vacuum on a **single** Postgres database container.

You can either run the crunchy-dba container as a single pod or include the container within a database pod.

The crunchy-dba container makes use of a Service Account to perform the startup of scheduled jobs. The Kube Job type is used to execute the scheduled jobs with a Restart policy of Never.

Environment Variables

The following environment variables control the actions of crunchy-dba:

- OSE_PROJECT - required, the OSE project name to log into
- JOB_HOST - required, the postgres container name the action will be taken against
- VAC_SCHEDULE - if set, this will start a vacuum job container. The setting value must be a valid cron expression as described below.
- BACKUP_SCHEDULE - if set, this will start a backup job container. The setting value must be a

valid cron expression as described below.

For a vacuum job, you are required to supply the following environment variables:

- PG_USER
- PG_PASSWORD
- JOB_HOST

Optional environment variables include:

- PG_DATABASE - defaults to postgres when not specified
- PG_PORT - defaults to 5432 when not specified
- VAC_ANALYZE - defaults to true when not specified
- VAC_FULL - defaults to true when not specified
- VAC_VERBOSE - defaults to true when not specified
- VAC_FREEZE - defaults to false when not specified
- VAC_TABLE - defaults to all tables when not specified, or you can set this value to indicate a single table to vacuum

CRON Expression Format

A cron expression represents a set of times, using 6 space-separated fields.

Table 1. Table Fields

Field name	Mandatory?	Allowed values	Allowed special characters
Seconds	Yes	0-59	* / , -
Minutes	Yes	0-59	* / , -
Hours	Yes	0-23	* / , -
Day of month	Yes	1-31	* / , - ?
Month	Yes	1-12 or JAN-DEC	* / , -
Day of week	Yes	0-6 or SUN-SAT	* / , - ?

Note: Month and Day-of-week field values are case insensitive. "SUN", "Sun", and "sun" are equally accepted.

Special Characters

Asterisk (*)

The asterisk indicates that the cron expression will match for all values of the field; e.g., using an asterisk in the 5th field (month) would indicate every month.

Slash (/)

Slashes are used to describe increments of ranges. For example 3-59/15 in the 1st field (minutes) would indicate the 3rd minute of the hour and every 15 minutes thereafter. The form "*\/..." is equivalent to the form "first-last/...", that is, an increment over the largest possible range of the field. The form "N/..." is accepted as meaning "N-MAX/...", that is, starting at N, use the increment until the end of that specific range. It does not wrap around.

Comma (,)

Commas are used to separate items of a list. For example, using "MON,WED,FRI" in the 5th field (day of week) would mean Mondays, Wednesdays and Fridays.

Hyphen (-)

Hyphens are used to define ranges. For example, 9-17 would indicate every hour between 9am and 5pm inclusive.

Question mark (?)

Question mark may be used instead of '*' for leaving either day-of-month or day-of-week blank.

Predefined schedules

You may use one of several pre-defined schedules in place of a cron expression.

Table 2. Table Predefined Schedules

Entry	Description	Equivalent To
@yearly (or @annually)	Run once a year, midnight, Jan. 1st	0 0 0 1 1 *
@monthly	Run once a month, midnight, first of month	0 0 0 1 * *
@weekly	Run once a week, midnight on Sunday	0 0 0 * * 0
@daily (or @midnight)	Run once a day, midnight	0 0 0 * * *
@hourly	Run once an hour, beginning of hour	0 0 * * * *

Intervals

You may also schedule a job to execute at fixed intervals. This is supported by formatting the cron spec like this:

```
@every <duration>
```

where "duration" is a string accepted by time.ParseDuration

(http://golang.org/pkg/time/#ParseDuration).

For example, "@every 1h30m10s" would indicate a schedule that activates every 1 hour, 30 minutes, 10 seconds.

Note: The interval does not take the job runtime into account. For example, if a job takes 3 minutes to run, and it is scheduled to run every 5 minutes, it will have only 2 minutes of idle time between each run.

Time zones

All interpretation and scheduling is done in the machines local time zone (as provided by the Go time package (http://www.golang.org/pkg/time). Be aware that jobs scheduled during daylight-savings leap-ahead transitions will not be run!

Legal Notices